Facial-Impression Judgments of Multimodal Large Language Models Predict Election Outcomes

Sirisa Cox[1] · James J. DiCarlo[2] · Gwangsu Kim[2]

[1]Commonwealth School [2]Massachusetts Institute of Technology Contact: kgspiano@mit.edu

SocialAgent: Second Workshop on LLMs for Social Reasoning and Simulation
NeurIPS 2026

MLLM judgments of which candidate looks more competent, based only on faces, predicted U.S. election winners above chance, and more accurately than the models' own vote choices.

Background

People infer traits from faces within a fraction of a second. These impressions are not necessarily accurate, but they predict consequential decisions, including hiring, sentencing, and voting. Competence judgments based only on candidates' faces predict election winners above chance (Todorov et al., 2005).

Recent work shows that multimodal large language models (MLLMs) reproduce human-like facial impressions. We asked whether MLLM competence judgments also predict real election outcomes.

Methods



"This is a study of first impressions from faces. Based only on appearance, if you had to pick one, which person looks more competent? Go with your immediate gut impression."

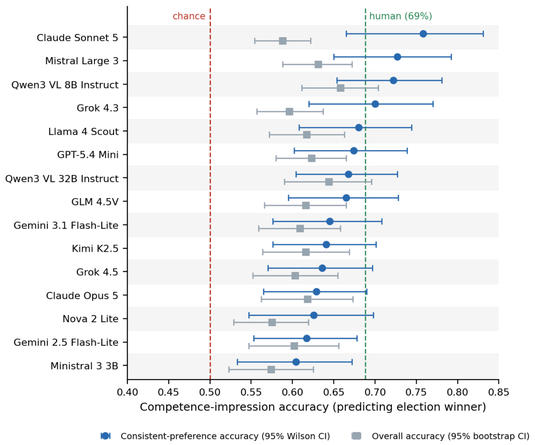
1. **Data.** 261 two-candidate U.S. gubernatorial and Senate races, 1995–2008 (Todorov-lab database), with human forced-choice competence judgments.
2. **Models.** 15 MLLMs with MMMU-Pro scores from 38% to 84%.
3. **Procedure.** Two-alternative forced choice, with each pair shown in both left/right orders over repeated trials.
4. **Scoring.** Overall accuracy includes every race. Consistent-preference accuracy includes only races in which the model chose the same candidate in both orders.

Competence judgments predict election outcomes

62% median overall accuracy (all races)

67% median consistent-preference accuracy (range 60–76%)

69% aggregate human competence judgments (234 races)



Circles: consistent-preference accuracy (Wilson 95% CI). Squares: overall accuracy (bootstrap 95% CI). Every model's confidence interval lies above chance.

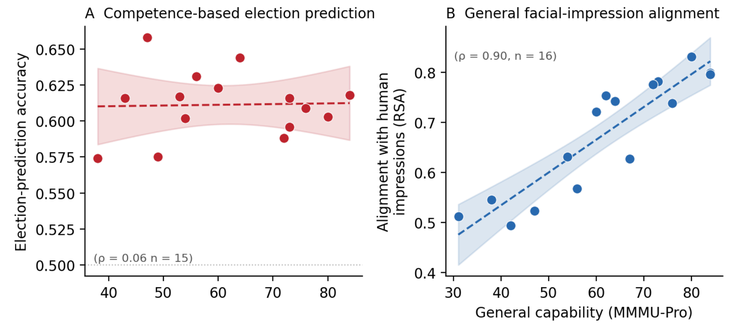
Comparison with other facial traits

Trait judged	Overall accuracy	Competence higher (N/15)
Competence	**0.61**	**–**
Familiarity	0.60	10/15 (n.s.)
Trustworthiness	0.56	12/15
Apparent age	0.54	14/15
Attractiveness	0.53	15/15
Dominance	0.51	15/15

Candidate recognition

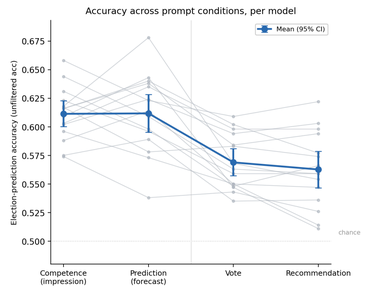
Most models correctly named fewer than 10% of candidates (highest: Kimi K2.5, 34%). Accuracy on races in which neither candidate was named was 61.6%, compared with 61.5% on all races. Accuracy was also above chance on open (non-incumbent) races.

Relationship to model capability



(A) No significant relationship between election accuracy and MMMU-Pro (ρ = 0.06, p = 0.82, n = 15; underpowered). (B) Alignment with human facial impressions increases with MMMU-Pro (ρ = 0.90, n = 16).

Impressions predict better than choices



Compared prompt	Competence higher	Diff. (pts)
Vote	15/15	+4.2
Recommend-ation	14/15	+4.9
Forecast	7/15	+0.0

Competence judgments were more accurate than answers to "which person would you vote for?" and "which person should I vote for?", and about as accurate as forecasts of the winner.

Conclusions

- As with human judgments, MLLM competence judgments from faces predict U.S. election outcomes above chance.
- These judgments were more predictive than the models' vote choices, so evaluations based only on decision-framed prompts may not detect this bias.
- Limitations: a single historical U.S. dataset, and contamination cannot be fully ruled out. A preregistered test on a future election would address the latter.